

# atJIT: a just-in-time autotuning compiler for C++

Kavon Farvardin • 23.07.2018

*Host: Hal Finkel*



# Just-in-time Compilation

Not a new idea:

1. Provide access to the compiler *while* executing the program.
2. The executable (re)compiles part of itself on-demand.

Why?

**Runtime information can significantly improve the quality of compiled code.**

**For computations like those at Argonne, this is very useful.**

```
int slowMul (int a, int b) {  
    int iters = b, answer = 0;  
    if (b < 0) {  
        // if b == -1*c  
        // then a * b == (a * -1) * c  
        a = -a;  
        iters = -b;  
    }  
    for (; iters > 0; --iters)  
        answer += a;  
    return answer;  
}
```

An implementation  
of integer multiplication  
that only uses addition &  
negation.

```

define i32 @slowMul(i32, i32) {
    %3 = icmp slt i32 %1, 0
    %4 = sub nsw i32 0, %1
    %5 = select i1 %3, i32 %4, i32 %1
    %6 = icmp sgt i32 %5, 0
    br i1 %6, label %7, label %16

; <label>:7:
    %8 = sub nsw i32 0, %0
    %9 = select i1 %3, i32 %8, i32 %0
    %10 = xor i32 %5, -1
    %11 = icmp sgt i32 %10, -2
    %12 = select i1 %11, i32 %10, i32 -2
    %13 = add i32 %5, %12
    %14 = add i32 %13, 2
    %15 = mul i32 %9, %14
    br label %16

; <label>:16:
    %17 = phi i32 [ 0, %2 ], [ %15, %7 ]
    ret i32 %17
}

```

A regular compiler can manage to optimize out the loop, but it's still inefficient!

```
const int b = // some dynamic value
for (int a = 0; a < 9999999; ++a) {
    ... = slowMul(a, b);
    ...
}
```



At start of loop,  
we observed b = 3

```
define i32 @slowMul(i32) {
    %2 = mul i32 %0, 3
    ret i32 %2
}
```

Let's dynamically compile  
a specialized version!

We got lucky and  
produced the best  
possible version!

```

-faggressive-loop-optimizations -falign-functions=<...> -falign-jumps=<...> -falign-labels=<...> -falign-loops=<...> -fassociative-math
-fauto-profile -fauto-profile[=path] -fauto-inc-dec -fbranch-probabilities -fbranch-target-load-optimize -fbranch-target-load-optimize2
-fbtr-bb-exclusive -fcaller-saves -fcombine-stack-adjustments -fconserve-stack -fcompare-elim -fcprop-registers -fcrossjumping -fcse-follow-jumps
-fcse-skip-blocks -fcx-fortran-rules -fcx-limited-range -fddata-sections -fdce -fdelayed-branch -fdelete-null-pointer-checks -fdevirtualize
-fdevirtualize-speculatively -fdevirtualize-at-ltrans -fdse -fearly-inlining -fipa-sra -fexpensive-optimizations -ffat-lto-objects -ffast-math
-ffinite-math-only -ffloat-store -fexcess-precision=style -fforward-propagate -ffp-contract=style -ffunction-sections -fgcse -fgcse-after-reload
-fgcse-las -fgcse-lm -fgraphite-identity -fgcse-sm -fhoist-adjacent-loads -fif-conversion -fif-conversion2 -findirect-inlining
-finline-functions -finline-functions-called-once -finline-limit=n -finline-small-functions -fipa-cp -fipa-cp-clone -fipa-bit-cp -fipa-vrp
-fipa-pta -fipa-profile -fipa-pure-const -fipa-reference -fipa-icf -fira-algorithm=algorithm -fira-region=region -fira-hoist-pressure
-fira-loop-pressure -fno-ira-share-save-slots -fno-ira-share-spill-slots -fisolate-erroneous-paths-dereference -fisolate-erroneous-paths-attribute
-fivopts -fkeep-inline-functions -fkeep-static-functions -fkeep-static-consts -flimit-function-alignment -flive-range-shrinkage -floop-block
-floop-interchange -floop-strip-mine -floop-unroll-and-jam -floop-nest-optimize -floop-parallelize-all -flra-remat -flto -flto-compression-level
-flto-partition=alg -fmerge-all-constants -fmerge-constants -fmodulo-sched -fmodulo-sched-allow-regmoves -fmove-loop-invariants
-fno-branch-count-reg -fno-defer-pop -fno-fp-int-builtin-inexact -fno-function-cse -fno-guess-branch-probability -fno-inline -fno-math-errno
-fno-peephole -fno-peephole2 -fno-printf-return-value -fno-sched-interblock -fno-sched-spec -fno-signed-zeros -fno-toplevel-reorder
-fno-trapping-math -fno-zero-initialized-in-bss -fomit-frame-pointer -foptimize-sibling-calls -fpartial-inlining -fpeel-loops
-fpredictive-commoning -fprefetch-loop-arrays -fprofile-correction -fprofile-use -fprofile-use=path -fprofile-values -fprofile-reorder-functions
-freciprocal-math -free -frename-registers -freorder-blocks -freorder-blocks-algorithm=algorithm -freorder-blocks-and-partition
-freorder-functions -frerun-cse-after-loop -freschedule-modulo-scheduled-loops -frounding-math -fsave-optimization-record -fsched2-use-superblocks
-fsched-pressure -fsched-spec-load -fsched-spec-load-dangerous -fsched-stalled-insns-dep[=n] -fsched-stalled-insns[=n] -fsched-group-heuristic
-fsched-critical-path-heuristic -fsched-spec-insn-heuristic -fsched-rank-heuristic -fsched-last-insn-heuristic -fsched-dep-count-heuristic
-fschedule-fusion -fschedule-insns -fschedule-insns2 -fsection-anchors -fselective-scheduling -fselective-scheduling2 -fsel-sched-pipelining
-fsel-sched-pipelining-outer-loops -fsemantic-interposition -fshrink-wrap -fshrink-wrap-separate -fsignaling-nans -fsingle-precision-constant
-fsplit-ivs-in-unroller -fsplit-loops -fsplit-paths -fsplit-wide-types -fssa-backprop -fssa-phiopt -fstdarg-opt -fstore-merging -fstrict-aliasing
-fthread-jumps -ftracer -ftree-bit-ccp -ftree-builtin-call-dce -ftree-ccp -ftree-ch -ftree-coalesce-vars -ftree-copy-prop -ftree-dce
-ftree-dominator-opts -ftree-dse -ftree-forwprop -ftree-fre -fcode-hoisting -ftree-loop-if-convert -ftree-loop-im -ftree-phi-prop
-ftree-loop-distribution -ftree-loop-distribute-patterns -ftree-loop-ivcanon -ftree-loop-linear -ftree-loop-optimize -ftree-loop-vectorize
-ftree-parallelize-loops=n -ftree-pre -ftree-partial-pre -ftree-pta -ftree-reassoc -ftree-sink -ftree-slsr -ftree-sra -ftree-switch-conversion
-ftree-tail-merge -ftree-ter -ftree-vectorize -ftree-vrp -funconstrained-commons -funit-at-a-time -funroll-all-loops -funroll-loops
-funsafe-math-optimizations -funswitch-loops -fipa-ra -fvariable-expansion-in-unroller -fvect-cost-model -fvpt -fweb -fwhole-program -fwpa
-fuse-linker-plugin --param name=value -O -O0 -O1 -O2 -O3 -Os -Ofast -Og

```



# Online autotuning

Objective: Easy-to-use automatic performance tuning via JIT compilation.

atJIT serves as a vehicle for experimenting with autotuning techniques.

Primary “knobs” we are focussing on:

1. Compiler Optimizations
  - a. Loop transforms
  - b. Inlining
  - c. Pass ordering/choices.
2. Algorithmic Choices

# Example Kernel

```
// multiply square matrices
template <typename T>
T** MatMul(const int DIM, T** aMatrix, T** bMatrix) {
    T** product = calloc_mat<T>(DIM);
    for (int row = 0; row < DIM; row++) {
        for (int col = 0; col < DIM; col++) {
            for (int inner = 0; inner < DIM; inner++) {
                product[row][col] += aMatrix[row][inner] * bMatrix[inner][col];
            }
        }
    }
    return product;
}
```



# Programmer's Interface

```
auto const &OptimizedFun = AT.reoptimize(MatMul<ElmTy>,  
    DIM, _1, _2, // arg 1 is fixed to DIM  
    tuner_kind(AT_Random));
```

```
ElmTy** ans = OptimizedFun(aMatrix, bMatrix);
```

---

Compile with a simple, drop-in replacement of of CXX:

```
➤ ../install/bin/easycc -O2 matmul.cpp -o matmul
```

# Automatic Tuning Process


avg time: 662929 ns, error: 0.401664%, ( $s^2$ : 3.10198e+07, s: 5954.09, SEM: 2662.75) from 5 measurements

```
{  
  inlining threshold := 225  
  loop #3 := <>  
  loop #2 := <>  
  loop #0 := <>  
  loop #1 := <>  
}
```

avg time: 445898 ns, error: 0.998717%, ( $s^2$ : 2.54827e+09, s: 50579.3, SEM: 4453.26) from 129 measurements

```
{  
  inlining threshold := 976  
  loop #3 := <llvm.loop.unroll.full: 1, llvm.loop.licm_versioning.disable: 0, llvm.loop.distribute.enable: 1, >  
  loop #2 := <llvm.loop.unroll.full: 1, llvm.loop.vectorize.enable: 1, llvm.loop.vectorize.width: 8, >  
  loop #0 := <llvm.loop.unroll.count: 32, llvm.loop.interleave.count: 1, llvm.loop.distribute.enable: 0, >  
  loop #1 := <llvm.loop.unroll.full: 1, llvm.loop.licm_versioning.disable: 1, llvm.loop.distribute.enable: 0, >  
}
```

# Matrix Multiply Random Tuning Results



Up to 1.48x Speedup, aka  
32.7% improvement!

Configuration Name



# Challenges

- Running-time normalization
- Parameter optimization
- Parallelizing JIT compilation
- ... and certainly more!

```
void performWork (dataset, iterations) { ... }
```

```
void main () {  
  ATDriver AT;
```

```
  auto func1 = AT.reoptimize(performWork, ...);  
  func1(dataset, 100);
```

```
  auto func2 = AT.reoptimize(performWork, ...);  
  func2(dataset, 500);  
  ...
```



# Conclusion

- I'm still here for 1¼ more months. My TODOs include:
  - Bayesian Optimization
  - Simulated Annealing
  - Algorithmic Tuning
- Special thanks to
  - Hal Finkel & Michael Kruse (Argonne)
  - John Reppy (UChicago)
  - Serge Guelton & Juan Manuel Martinez Caamaño  
(developed open-source infrastructure we started from)



Code is available on GitHub:

<https://github.com/kavon/atJIT>